

Non-Intrusive Parallel-in-Time Solvers for Challenging Problems

Research in Industry Projects for Student

Tran Duy Anh Le¹, Bryan Li², Margaret Luo³,
Daniel Agraz Vallejo⁴

Academic Mentor: Jean-Michel Maldague⁵

Industry Mentors: Robert Falgout⁶, Rui Peng Li⁶, Wayne Mitchell⁶, Daniel Osei-Kuffuor⁶

¹University of Rochester, ²University of California, Berkeley, ³University of California, San Diego, ⁴CETYS University

⁵University of California, Los Angeles, ⁶Lawrence Livermore National Laboratory



August 20, 2024



Lawrence Livermore National Laboratory (LLNL)



Lawrence Livermore National Laboratory (LLNL)



- LLNL is a federally funded research institution located in Livermore, California established in 1952.
- Home of the World's fastest supercomputers including Sierra and El Capitan.
- Over 8,000 people works at LLNL.

LLNL's Center for Applied Scientific Computing (CASC)



- LLNL's branch that hosts 11 research groups that gathers various multidisciplinary teams to conduct scientific research.

| | | | | | |
|---|--|---|--------------------------------|---|---|
|  | Computational Physics Group |  | Data Science & Analytics Group |  | High Performance Computing Group |
|  | Informatics Group |  | Machine Intelligence Group |  | Mathematical Algorithms & Computing Group |
|  | Numerical Analysis & Simulations Group |  | Parallel Systems Group |  | Research Software Engineering Group |
|  | Scientific Computing Group |  | UQ Optimization Group | | |

The need for parallelism

- Serial algorithms are bottlenecked on single-thread performance.
- Parallel algorithms take advantage of core count, which is growing exponentially.
 - Modern supercomputers (such as LLNL's Sierra) have core counts $\sim 10^6\text{--}10^7$.

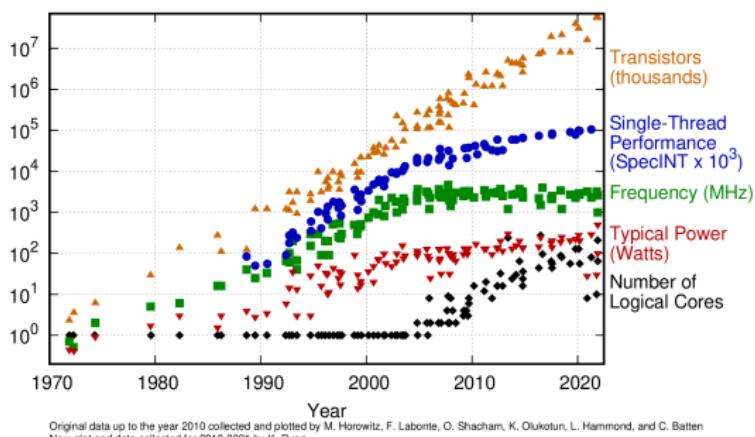


Figure: Stagnation of processor speed

Parallel-in-time algorithms

- Our target is parallel algorithms to simulate **evolutionary partial differential equations**.
- These model many physical phenomena that LLNL is interested in:
 - Heat transfer
 - Wave propagation
 - Quantum evolution
- Serial-in-time algorithms require many timesteps, which cannot be evaluated in parallel.
- Parallel-in-time algorithms evaluate timesteps in parallel, taking advantage of more cores to provide speedup.
 - This is the focus of our project!

Intrusive vs. Non-Intrusive Algorithms

Intrusive:

- Significant effort and additional code required to apply to new problems.
- Approach usually very heavily **problem-dependent**.

Intrusive vs. Non-Intrusive Algorithms

Intrusive:

- Significant effort and additional code required to apply to new problems.
- Approach usually very heavily **problem-dependent**.

Non-Intrusive:

- Little effort and additional code required to apply to new problems.
- Approach/algorithm **generalizes** to new problems.

Discretization using finite difference

- To deal with PDEs using a computer, we must first discretize them.
- For example, consider the 1D *steady-state diffusion* (Poisson) equation

$$-u_{xx} = f.$$

on the domain $[0, 1]$.

Discretization using finite difference

- To deal with PDEs using a computer, we must first discretize them.
- For example, consider the 1D *steady-state diffusion* (Poisson) equation

$$-u_{xx} = f.$$

on the domain $[0, 1]$.

- Divide domain up into N equal pieces of length $\delta x = 1/N$.
- Set nodes $x_i = i\delta x$ for $i = 0, \dots, N$, and let $u_i \approx u(x_i)$, $f_i = f(x_i)$.

Discretization using finite difference

- To deal with PDEs using a computer, we must first discretize them.
- For example, consider the 1D *steady-state diffusion* (Poisson) equation

$$-u_{xx} = f.$$

on the domain $[0, 1]$.

- Divide domain up into N equal pieces of length $\delta x = 1/N$.
- Set nodes $x_i = i\delta x$ for $i = 0, \dots, N$, and let $u_i \approx u(x_i)$, $f_i = f(x_i)$.
- We can use the *finite difference approximation*

$$-u_{xx}(x_i) \approx \frac{-u_{i-1} + 2u_i - u_{i+1}}{\delta x^2}.$$

Discretization using finite difference

- Using this approximation, for $i = 1, \dots, N - 1$, we obtain

$$\frac{-u_{i-1} + 2u_i - u_{i+1}}{\delta x^2} = f_i.$$

- Assuming Dirichlet boundary conditions $u_0 = u_N = 0$, we can write this in a linear system $A\mathbf{u} = \mathbf{f}$:

$$\frac{1}{\delta x^2} \begin{pmatrix} 2 & -1 & 0 & \cdots & 0 \\ -1 & 2 & -1 & & \vdots \\ 0 & -1 & \ddots & & 0 \\ \vdots & & & & -1 \\ 0 & \cdots & 0 & -1 & 2 \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \\ \vdots \\ u_{N-1} \end{pmatrix} = \begin{pmatrix} f_1 \\ f_2 \\ \vdots \\ f_{N-1} \end{pmatrix}.$$

Relaxation methods

- Finite difference discretization produces a linear system we need to solve.
- Linear systems formed from discretization can be very large.
- Directly solving (e.g. Gaussian elimination) results in $O(n^3)$ runtime: too slow!
- **Relaxation methods**, which are iterative methods for solving linear systems, typically with $O(n^2)$ or better runtime per iteration, are used instead.

Relaxation methods and Multigrid

Relaxation methods are better at correcting oscillatory errors, but are worse at correcting low-frequency errors.



Figure: Fine Grid Iteration 0

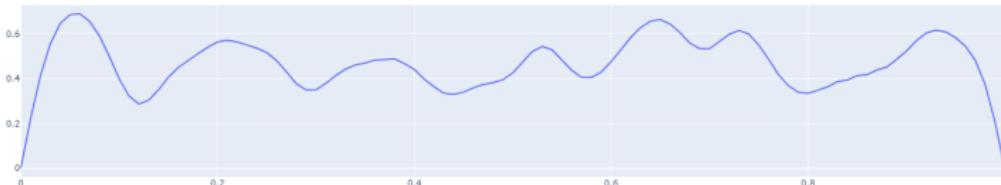


Figure: Fine Grid Iteration 5

Multigrid: the V-cycle

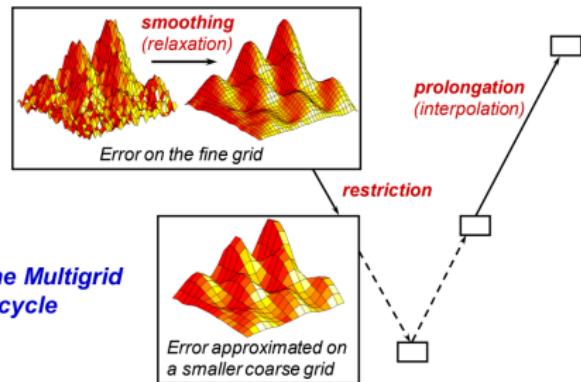


Figure: Multigrid V-cycle

- Multigrid attacks the weakness of relaxation methods by going through a hierarchy of grid sizes.

Multigrid: the V-cycle

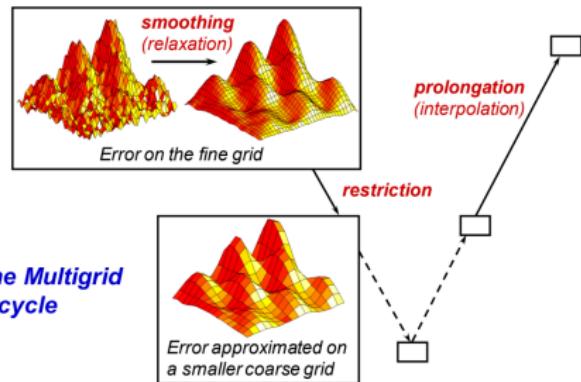


Figure: Multigrid V-cycle

- Multigrid attacks the weakness of relaxation methods by going through a hierarchy of grid sizes.
- Low-frequency (smooth) errors on the fine grid turn into high-frequency (oscillatory) error on the coarser grid and are corrected more effectively through relaxation.

Relaxation methods and Multigrid

Relaxation methods are better at correcting oscillatory errors.



Figure: Fine Grid Iteration 0



Figure: Fine Grid Iteration 5



Figure: Coarse Grid Iteration 0



Figure: Coarse Grid Iteration 5

Multigrid reduction in time (MGRIT)

MGRIT algorithm is a parallel-in-time approach for solving time dependent problems that is designed to be non-intrusive.

MGRIT

Consider the diffusion equation $u_t - u_{xx} = f$ with the following discretization

$$u_t \approx \frac{u_i^{j+1} - u_i^j}{\delta t}, \quad -u_{xx} \approx \frac{-u_{i-1}^j + 2u_i^j - u_{i+1}^j}{\delta x^2}$$

where $u_i^j \approx u(x_i, t_j)$.

Using Forward Euler, we get the update rule

$$\mathbf{u}^j = \Phi \mathbf{u}^{j-1} + \delta t \mathbf{f}^{j-1} \quad \text{for } j = 1, 2, \dots, n \text{ and } \mathbf{u}^0 = \mathbf{g}^0,$$

where $\mathbf{u}^j \approx [u_0^j \quad u_1^j \quad \dots \quad u_N^j]^T$ and Φ has the stencil $[\beta \quad 1 - 2\beta \quad \beta]$, where $\beta = \delta t / \delta x^2 \leq 1/2$ by stability condition.

MGRIT: the coarse-grid operator

Using the partition below, we let $t_j = j\delta t$ for the fine time grid and $T_j = mj\delta t$ for the coarse time grid with coarsening factor m .

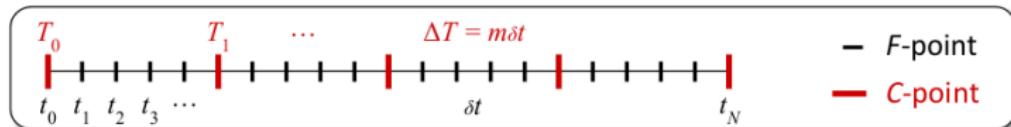


Figure: Coarse and fine time grids

Then to solve each coarse time step is to solve the following system

$$T(m)\mathbf{u} = \begin{pmatrix} I & & & \\ -\Phi^m & I & & \\ & -\Phi^m & I & \\ & & \ddots & \ddots & \\ & & & -\Phi^m & I \end{pmatrix} \begin{pmatrix} u_0 \\ u_m \\ u_{2m} \\ \vdots \\ u_N \end{pmatrix} = \mathbf{f}.$$

Approximating the coarse-grid operator

However, Φ^m is expensive to compute with, thus we ultimately want to **approximate Φ^m with a sparser operator Ψ .**

Approximating the coarse-grid operator

However, Φ^m is expensive to compute with, thus we ultimately want to **approximate Φ^m with a sparser operator Ψ .**

So we have the “target matrix” $T(m)$ which we want to approximate with some matrix A

$$T(m) = \begin{pmatrix} I & & & \\ -\Phi^m & I & & \\ & \ddots & \ddots & \\ & & -\Phi^m & I \end{pmatrix}, \quad A = \begin{pmatrix} I & & & \\ -\Psi & I & & \\ & \ddots & \ddots & \\ & & -\Psi & I \end{pmatrix}.$$

Approximating the coarse-grid operator (cont'd)

$$T(m) = \begin{pmatrix} I & & & \\ -\Phi^m & I & & \\ & \ddots & \ddots & \\ & & -\Phi^m & I \end{pmatrix}, \quad A = \begin{pmatrix} I & & & \\ -\Psi & I & & \\ & \ddots & \ddots & \\ & & -\Psi & I \end{pmatrix}$$

For a good approximation, we want A to be *spectrally equivalent* to $T(m)$, i.e. the following is bounded independent of problem size

$$\|I - AT^{-1}\|.$$

$$\begin{aligned} \|I - AT^{-1}\| &= \sup_x \frac{\|(I - AT^{-1})x\|}{\|x\|} \\ &= \sup_y \frac{\|(T - A)y\|}{\|Ty\|} \end{aligned}$$

Why Machine Learning?

We will utilize Machine Learning techniques for this project.

Why Machine Learning?

We will utilize Machine Learning techniques for this project.

- Current approaches of approximating Φ^m applicable to only certain types of PDEs.

Why Machine Learning?

We will utilize Machine Learning techniques for this project.

- Current approaches of approximating Φ^m applicable to only certain types of PDEs.
- We want a “non-intrusive” way to estimate coarse-grid operators.

Why Machine Learning?

We will utilize Machine Learning techniques for this project.

- Current approaches of approximating Φ^m applicable to only certain types of PDEs.
- We want a “non-intrusive” way to estimate coarse-grid operators.
- Machine learning is good at coming up at more general solutions.

Neural Network and PyMGRIT

- Construct a Neural Network whose output is spectrally equivalent to the target matrix

Neural Network and PyMGRIT

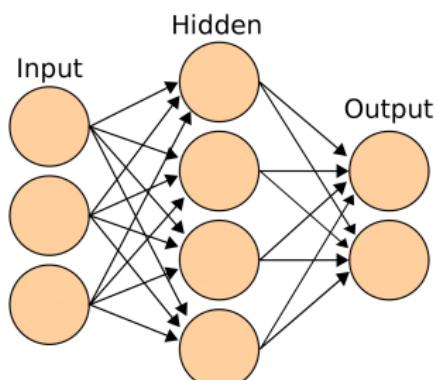
- Construct a Neural Network whose output is spectrally equivalent to the target matrix
- Evaluate the performance of the Neural Network model integrated into PyMGRIT

Neural Network Architecture

For 1D diffusion problem:

Input: $(\phi, m) \in \mathbb{R}^3 \times \mathbb{N}$.

Output: $\psi \in \mathbb{R}^3$.



ϕ is the “stencil” of Φ . For example, in our periodic problem,

$$L = \frac{1}{\delta x^2} \begin{bmatrix} 2 & -1 & & & -1 \\ -1 & 2 & -1 & & \\ & \ddots & \ddots & \ddots & \\ & & -1 & 2 & -1 \\ -1 & & & -1 & 2 \end{bmatrix},$$

and $\Phi = I - \delta t L$ has stencil

$\phi = [\beta \quad 1 - 2\beta \quad \beta]$ where $\beta = \delta t / \delta x^2$

Loss Function

- Goal: Find an A that's *spectrally equivalent* to T . In other words, minimize $\|I - AT^{-1}\|$.

Loss Function

- Goal: Find an A that's *spectrally equivalent* to T . In other words, minimize $\|I - AT^{-1}\|$.
- Recall, the ℓ^2 norm of a vector $v \in \mathbb{R}^d$ is $\|v\|_{\ell^2}^2 = v_1^2 + \cdots + v_d^2$.

Loss Function

- Goal: Find an A that's *spectrally equivalent* to T . In other words, minimize $\|I - AT^{-1}\|$.
- Recall, the ℓ^2 norm of a vector $v \in \mathbb{R}^d$ is $\|v\|_{\ell^2}^2 = v_1^2 + \cdots + v_d^2$.
- The norm of a matrix is defined as the induced 2-norm (or spectral norm)

$$\|A\| = \sup_{0 \neq v \in \mathbb{R}^d} \frac{\|Av\|_{\ell^2}}{\|v\|_{\ell^2}}.$$

Loss Function

- Goal: Find an A that's *spectrally equivalent* to T . In other words, minimize $\|I - AT^{-1}\|$.
- Recall, the ℓ^2 norm of a vector $v \in \mathbb{R}^d$ is $\|v\|_{\ell^2}^2 = v_1^2 + \cdots + v_d^2$.
- The norm of a matrix is defined as the induced 2-norm (or spectral norm)

$$\|A\| = \sup_{0 \neq v \in \mathbb{R}^d} \frac{\|Av\|_{\ell^2}}{\|v\|_{\ell^2}}.$$

- Equivalently, this is the largest magnitude of eigenvalue of A .

Loss Function 1

By choosing some vectors $\{y_k\}_{k=1}^K$, we can get the following loss function

Loss Function 1

By choosing some vectors $\{y_k\}_{k=1}^K$, we can get the following loss function

1st Loss function

$$\mathcal{L}_1 = \frac{\|Ty_k - Ay_k\|^2}{\|Ty_k\|^2} = \frac{\sum_r (T_r y_k - A_r y_k)^2}{\sum_r (T_r y_k)^2} \leq \sum_r \frac{(T_r y_k - A_r y_k)^2}{(T_r y_k)^2}$$

where T_r, A_r denote row r of T and A respectively.

Loss Function 1

By choosing some vectors $\{y_k\}_{k=1}^K$, we can get the following loss function

1st Loss function

$$\mathcal{L}_1 = \frac{\|Ty_k - Ay_k\|^2}{\|Ty_k\|^2} = \frac{\sum_r (T_r y_k - A_r y_k)^2}{\sum_r (T_r y_k)^2} \leq \sum_r \frac{(T_r y_k - A_r y_k)^2}{(T_r y_k)^2}$$

where T_r, A_r denote row r of T and A respectively.

By focusing on one row, and letting $y_k(i)$ be the restriction of y_k to the coarse time point i , we obtain a simplified expression

$$\mathcal{L}'_1 = \sum_{k=1}^K \frac{((\Phi^m)_r y_k - \Psi_r y_k(i))^2}{(1 - (\Phi^m)_r y_k(i))^2}.$$

Loss Function 2

Our second idea for a loss function is inspired by Ru Huang et al (2023).

Loss Function 2

Our second idea for a loss function is inspired by Ru Huang et al (2023). Let (λ_k, y_k) be the orthogonal eigenpairs of Φ^m , sorted such that $|\lambda_i| \geq |\lambda_{i+1}|$. Let $K \leq d$ be the number of testing vectors.

Loss Function 2

Our second idea for a loss function is inspired by Ru Huang et al (2023). Let (λ_k, y_k) be the orthogonal eigenpairs of Φ^m , sorted such that $|\lambda_i| \geq |\lambda_{i+1}|$. Let $K \leq d$ be the number of testing vectors. Then we get the following loss function

2nd Loss Function

$$\mathcal{L}_2 = \sum_{k=1}^K \|\Phi^m y_k - \Psi y_k\|_{\ell^2}^2.$$

Loss Function 3

Another loss function is developed similarly to the second one.

Loss Function 3

Another loss function is developed similarly to the second one.
Consider (λ_k, y_k) be orthogonal eigenpairs of Φ , sorted such that
 $|\lambda_i| \geq |\lambda_{i+1}|$. Let $K \leq d$ be the number of testing vectors.

Loss Function 3

Another loss function is developed similarly to the second one.

Consider (λ_k, y_k) be orthogonal eigenpairs of Φ , sorted such that $|\lambda_i| \geq |\lambda_{i+1}|$. Let $K \leq d$ be the number of testing vectors.

By expanding out the ℓ_2 norm of the vectors, if M is the number of block rows of $T(V)$, we have a new loss function.

3rd Loss Function

$$\mathcal{L}_{3,M} = \sum_{k=1}^K \frac{(M-1)\|(\Psi - \Phi^m)v_k\|_{\ell_2}^2}{1 + (M-1)|1 - \lambda_k^m|^2}$$

Currently, M is chosen to be a sufficiently large number or $M = \frac{N}{m} + 1$ where N is a large number and m is the coarsening factor.

Error analysis

- Consider $\{\lambda_k\}$ be the set of eigenvalues of Φ .

Error analysis

- Consider $\{\lambda_k\}$ be the set of eigenvalues of Φ .
- For diffusion-dominated equations, there are very few $|\lambda_k| \approx 1$ and many $|\lambda_k| \ll 1$.

Error analysis

- Consider $\{\lambda_k\}$ be the set of eigenvalues of Φ .
- For diffusion-dominated equations, there are very few $|\lambda_k| \approx 1$ and many $|\lambda_k| \ll 1$.
- For example, for the equation $u_t = u_{xx}$, if we consider N gridpoints, then the eigenvalues are given by

$$\lambda_k = 1 - 2\beta \left(1 - \cos \left(\frac{k\pi}{N+1} \right) \right), \quad 1 \leq k \leq N$$

where $\beta = \delta t / \delta x^2 \leq 1/2$.

Error analysis

- Consider $\{\lambda_k\}$ be the set of eigenvalues of Φ .
- For diffusion-dominated equations, there are very few $|\lambda_k| \approx 1$ and many $|\lambda_k| \ll 1$.
- For example, for the equation $u_t = u_{xx}$, if we consider N gridpoints, then the eigenvalues are given by

$$\lambda_k = 1 - 2\beta \left(1 - \cos \left(\frac{k\pi}{N+1} \right) \right), \quad 1 \leq k \leq N$$

where $\beta = \delta t / \delta x^2 \leq 1/2$.

- This makes approximation for Ψ easier as smaller eigenvalues' decay faster.

Error analysis

- Consider $\{\lambda_k\}$ be the set of eigenvalues of Φ .
- For diffusion-dominated equations, there are very few $|\lambda_k| \approx 1$ and many $|\lambda_k| \ll 1$.
- For example, for the equation $u_t = u_{xx}$, if we consider N gridpoints, then the eigenvalues are given by

$$\lambda_k = 1 - 2\beta \left(1 - \cos \left(\frac{k\pi}{N+1} \right) \right), \quad 1 \leq k \leq N$$

where $\beta = \delta t / \delta x^2 \leq 1/2$.

- This makes approximation for Ψ easier as smaller eigenvalues' decay faster.
- However, the distribution of $|\lambda_k|$ varies between different types of PDEs (hyperbolic, elliptic, etc).

Loss Functions we are comparing

- $\mathcal{L}_1 = \sum_{k=1}^K \frac{((\Phi^m)_r y_k - \Psi_r y_k(i))^2}{(1 - (\Phi^m)_r y_k(i))^2 + \epsilon}$ where y_k are random vectors with an additional constant vector of all 1's, and ϵ is added to avoid dividing by zero.
- $\mathcal{L}_1 = \sum_{k=1}^K \frac{((\Phi^m)_r y_k - \Psi_r y_k(i))^2}{(1 - (\Phi^m)_r y_k(i))^2 + \epsilon}$ where y_k are largest eigenvectors of Φ .
- $\mathcal{L}_2 = \sum_{k=1}^K \|\Phi^m y_k - \Psi y_k\|_{\ell^2}^2$ where y_k are largest eigenvectors of Φ .
- $\mathcal{L}_{3,M} = \sum_{k=1}^K \frac{(M-1)\|\Phi^m y_k - \Psi y_k\|_{\ell^2}^2}{(1 + (M-1)|\lambda_k^m - 1|)^2}$ where (λ_k, y_k) are largest pairs of eigenvalues and eigenvectors of Φ
- $\mathcal{L}_{3,\text{Hybrid}}$, which has a similar structure to $\mathcal{L}_{3,M}$ but we update M during training.

Loss Functions we are comparing (cont'd)

We also consider modifications of the preceding loss functions, including:

- Using eigenvectors instead of random vectors for \mathcal{L}_1 .
- Adding the constant vector of all 1s to the test vector list for \mathcal{L}_2 , \mathcal{L}_3 and \mathcal{L}_1 with eigenvectors.

Loss Function Landscapes

Consider the problem set up with

$$\delta x = \frac{1}{16}, \quad \delta t = \frac{1}{4096}, \quad \beta = \frac{1}{16} \quad m = 4$$

The stencil for $\Phi = [0.0625 \quad 0.8750 \quad 0.0625]$, and our input to the loss function is $[0.0625, 0.8750, 0.0625, 4.0000]$.

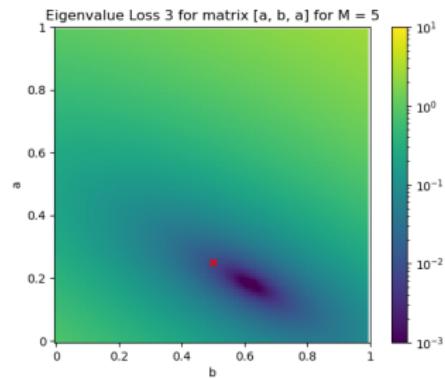
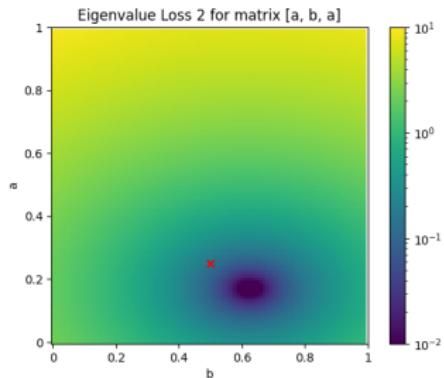
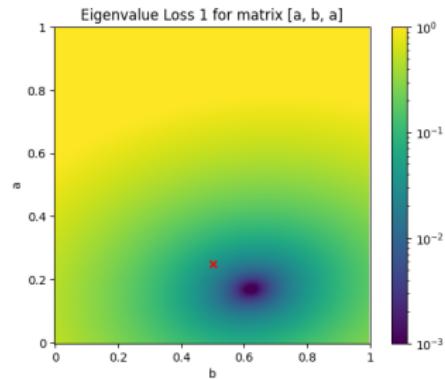
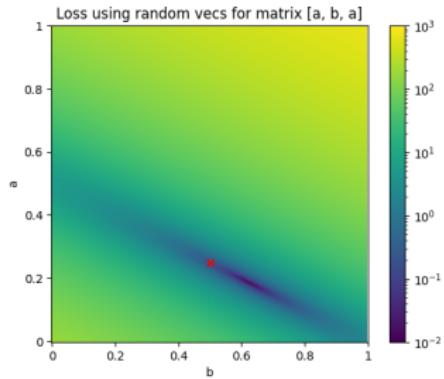
The re-discretization stencil is $[0.2500 \quad 0.5000 \quad 0.2500]$

If we restrict the approximation Ψ to have stencil of the form $[a \quad b \quad a]$, we can plot the loss at (a, b) for each loss function.

This allows us to observe the loss landscape for each loss function, as well as where the re-discretization stencil, which we know converges well in PyMGRIT, lies.

Loss Function Landscapes (cont'd)

✗ is the re-discretization stencil.



Loss Function Landscapes (cont'd)

\mathcal{L}_1 with eigenvectors vs eigenvectors and constant vector.

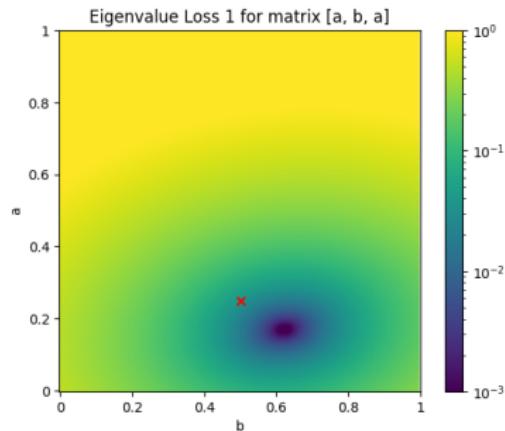


Figure: \mathcal{L}_1 with evecs

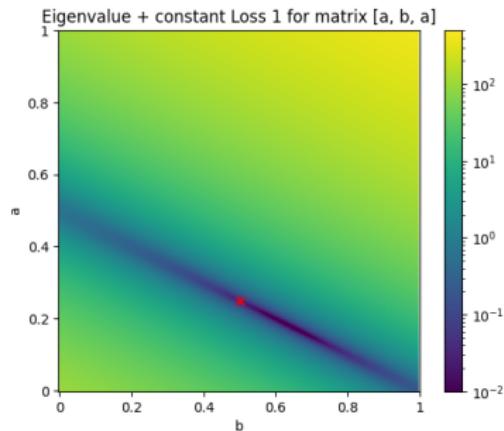


Figure: \mathcal{L}_1 with evecs and the constant vector

Loss Function Landscapes (cont'd)

\mathcal{L}_2 with eigenvectors vs eigenvectors and constant vector.

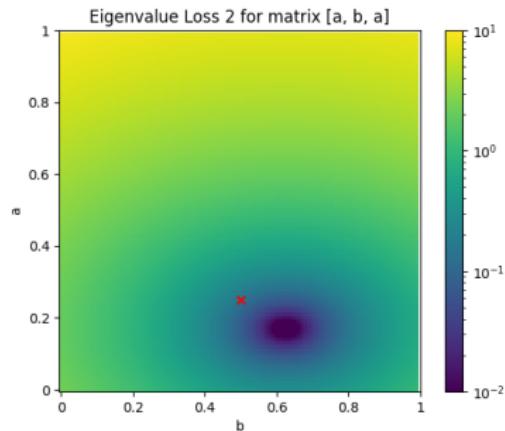


Figure: \mathcal{L}_2 with evecs

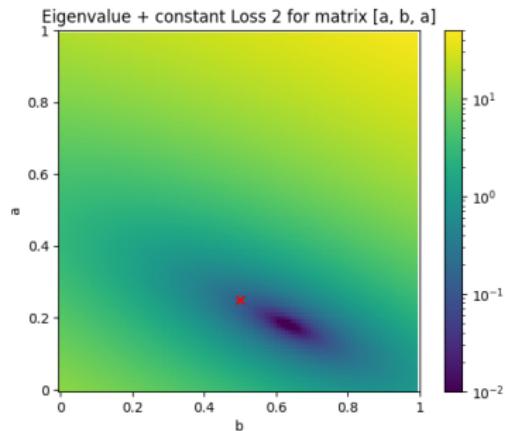


Figure: \mathcal{L}_2 with evecs and the constant vector

Loss Function Landscapes (cont'd)

$\mathcal{L}_{3,5}$ with eigenvectors vs eigenvectors and constant vector.

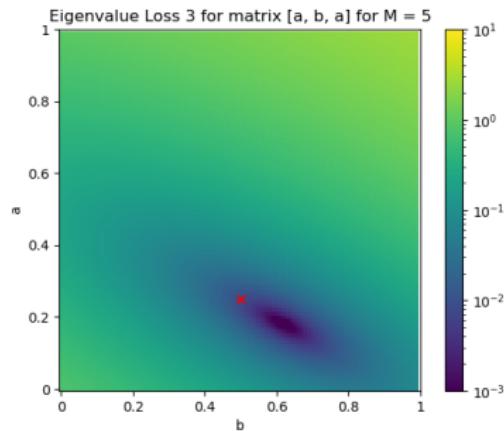


Figure: $\mathcal{L}_{3,5}$ with evecs

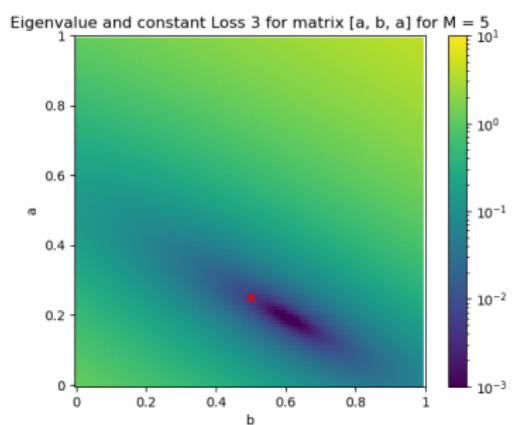


Figure: $\mathcal{L}_{3,5}$ with evecs and the constant vector

Performance of optimal stencil in PyMGRIT

$$\delta x = \frac{1}{16}, \quad \delta t = \frac{1}{4096}, \quad m = 4$$

| Loss Function | Optimal Stencil | Residual | Conv. Factor |
|-------------------------------|------------------------|----------|--------------|
| \mathcal{L}_1 Random | [0.1814 0.6324 0.1814] | 6.60e-5 | 4.04e-1 |
| \mathcal{L}_1 EVecs | [0.1700 0.6221 0.1700] | 2.13e-2 | 9.40e-1 |
| \mathcal{L}_1 EVecs, const. | [0.1824 0.6341 0.1824] | 6.60e-6 | 2.82e-1 |
| \mathcal{L}_2 | [0.1705 0.6223 0.1705] | 2.11e-2 | 9.59e-1 |
| \mathcal{L}_2 const. | [0.1797 0.6346 0.1797] | 1.55e-4 | 4.17e-1 |
| $\mathcal{L}_{3,5}$ | [0.1891 0.6071 0.1891] | 6.82e-2 | 9.03e-1 |
| $\mathcal{L}_{3,5}$, const. | [0.1977 0.5945 0.1977] | 3.32e-2 | 7.62e-1 |
| $\mathcal{L}_{3,20}$ | [0.2028 0.5895 0.2028] | 1.63e-2 | 4.48e-1 |
| $\mathcal{L}_{3,20}$, const. | [0.2061 0.5840 0.2061] | 2.87e-3 | 3.47e-1 |
| Re-disc | [0.2500 0.5000 0.2500] | 1.73e-8 | 9.93e-2 |

Neural Network Testing Setup

Training set: $\beta \in \{1/8, 1/12, 1/16, 1/24\}$, $m = \{1, 2, 3, 4\}$ (4 stencils).

Test problem: $\beta = 1/10$, $\delta x = 1/16$, $\delta t = 1/2560$. In PyMGRIT,
 $n_t = 2561$, $n_x = 17$.

We use a feedforward neural network with 3 hidden layers of 50 neurons each with the LeakyReLU activation function.

Neural Network Results ($m = 2$)

| Loss Function | Output Stencil | Output Loss | Redisc. Loss |
|--|---------------------|-------------|--------------|
| \mathcal{L}_1 Random | [0.131 0.743 0.131] | 0.177 | 0.143 |
| \mathcal{L}_1 Evecs | [0.127 0.726 0.121] | 6.75e-3 | 6.15e-3 |
| \mathcal{L}_1 Evecs, const. | [0.159 0.693 0.150] | 1.21e-3 | 4.83e-3 |
| \mathcal{L}_2 | [0.128 0.733 0.120] | 4.03e-2 | 3.48e-2 |
| \mathcal{L}_2 const. | [0.130 0.738 0.125] | 4.15e-2 | 3.48e-2 |
| $\mathcal{L}_{3,5}$ | [0.209 0.509 0.255] | 9.64e-3 | 8.07e-4 |
| $\mathcal{L}_{3,5}$ const. | [0.159 0.664 0.165] | 1.11e-3 | 3.62e-3 |
| $\mathcal{L}_{3,10}$ | [0.223 0.505 0.256] | 7.26e-3 | 5.67e-4 |
| $\mathcal{L}_{3,10}$ const. | [0.168 0.666 0.166] | 1.53e-3 | 2.76e-3 |
| $\mathcal{L}_{3,20}$ | [0.239 0.485 0.266] | 5.97e-3 | 3.40e-4 |
| $\mathcal{L}_{3,20}$ const. | [0.196 0.619 0.199] | 5.97e-3 | 3.40e-4 |
| $\mathcal{L}_{3,\text{Hybrid}}$ | [0.246 0.473 0.275] | 5.14e-3 | 2.48e-4 |
| $\mathcal{L}_{3,\text{Hybrid}}$ const. | [0.243 0.420 0.334] | 5.05e-2 | 1.47e-3 |
| Re-disc. | [0.200 0.600 0.200] | — | — |

PyMGRIT Results ($m = 2$)

| Loss Function | Output Stencil | Residual | Conv. Factor |
|--|---------------------|----------|--------------|
| \mathcal{L}_1 Random | [0.131 0.743 0.131] | 2.28 | 5.74e-1 |
| \mathcal{L}_1 Evecs | [0.127 0.726 0.121] | 1.85e-2 | 9.42e-1 |
| \mathcal{L}_1 Evecs, const. | [0.159 0.693 0.150] | 2.07e-4 | 4.03e-1 |
| \mathcal{L}_2 | [0.128 0.733 0.120] | 1.13e-2 | 9.03e-1 |
| \mathcal{L}_2 const. | [0.130 0.738 0.125] | 1.69e-3 | 6.3e-1 |
| $\mathcal{L}_{3,5}$ | [0.209 0.509 0.255] | 9.25e-2 | 9.86e-1 |
| $\mathcal{L}_{3,5}$ const. | [0.159 0.664 0.165] | 6.40e-2 | 9.13e-1 |
| $\mathcal{L}_{3,10}$ | [0.223 0.505 0.256] | 7.84e-2 | 9.57e-1 |
| $\mathcal{L}_{3,10}$ const. | [0.168 0.666 0.166] | 2.19e-8 | 1.67e-1 |
| $\mathcal{L}_{3,20}$ | [0.239 0.485 0.266] | 5.08e-2 | 8.64e-1 |
| $\mathcal{L}_{3,20}$ const. | [0.196 0.619 0.199] | 1e12 | 1.75 |
| $\mathcal{L}_{3,\text{Hybrid}}$ | [0.246 0.473 0.275] | 1.43e-2 | 6.57e-1 |
| $\mathcal{L}_{3,\text{Hybrid}}$ const. | [0.243 0.420 0.334] | 1.54e-2 | 6.81e-1 |
| Re-disc. | [0.200 0.600 0.200] | 7.60e-8 | 6.77e-2 |

Further research recommendations

- 1D Diffusion only proof-of-concept – generalize to more types of PDEs, such as advection (which is harder for MGRIT).
- Currently, Ψ hard-coded to be tridiagonal. For larger coarsening factors, larger stencil may result in better convergence. Use second neural network to select nonzero entries of Ψ , following Ru Huang et als' paper.
- Use bootstrapping ideas:
 - Use the output to train the model.
 - Solve $A\mathbf{u} = \mathbf{0}$ to generate error vectors and update the training set with them.
- Work on periodic condition.

Acknowledgement



Figure: Lawrence Livermore National Laboratory Site Visit Photo

Q & A

- Thank you for listening!
- Questions?